



## Hackathon Setup guide

Leverage featherless.ai - a serverless AI inference platform with unlimited access to thousands of open-source AI models. Focus on building amazing applications while we handle the infrastructure.

# Get Started in 3 Steps



## Sign up for Feather Premium

- Sign up at [featherless.ai](https://featherless.ai),
- go to the Pricing page
- Select Feather Premium and fill in coupon code **ENVIRO**

The screenshot shows the featherless.ai website's Pricing page. At the top, there's a navigation bar with links for Models, Resources, Pricing, Chat, and Status. A search bar and a user profile icon are also present.

In the center, there's a section titled "Upgrade your plan" with a back arrow and the text "Upgrade your plan". Below this, there are two sections: "Billing Information" and "Payment Information".

**Billing Information:** Fields for "Full name", "Country or region" (set to Belgium), and "Address".

**Payment Information:** Buttons for "Card" (highlighted in green), "Google Pay", "Bank", and "US\$5 back". Below these are links for "Secure, fast checkout with Link" and payment method icons for "Card", "Visa", and "Mastercard".

**Feather Premium Plan:** A yellow button labeled "\$25.00 /month" with a "Change" link. To the right, a list of features includes:

- Access to DeepSeek, Kimi-K2 and GLM 4.6
- Access any model - no limit on size!
- Up to 4 concurrent connections
- Up to 32K context

**Charging Summary:** A table showing charges:

You will be charged the following:	
1 × Feather Premium (at \$25.00 / month)	\$25.00
Amount due now	\$25.00

A red arrow points from the "Amount due now" row to the "Got a coupon or referral code?" input field.

**Coupon Input:** A red-bordered input field with placeholder text "Enter promotion code here" and a "Apply" button.

By providing your card information, you allow Featherless AI to charge your card for future purchases.

# Get Started in 3 Steps



## Get your API key

- After subscribing, go to the top right and click on API Keys
- Here you can create an API key, keep this API key handy, you'll need it to access the service

The screenshot shows the featherless.ai dashboard. On the left, there's a sidebar with 'FEATHERLESS' and 'ACCOUNT' sections. Under 'FEATHERLESS', 'API Keys' is selected and highlighted with a red box. Under 'ACCOUNT', there are links for 'Profile', 'Change Password', 'Subscription', 'Payment Methods', and 'Referral Program'. At the top, there are navigation links for 'Models', 'Resources', 'Pricing', 'Chat', and 'Status'. A green dot next to 'Status' indicates it's active. In the center, the 'API Keys' section displays two keys: 'aifu' and 'litellm'. Each key has columns for 'Key', 'Created', and 'Last Used'. To the right, a user profile for 'Darin Verheijke' is shown with a blue profile picture. The 'Account & Billing' section is expanded, showing 'API Keys' (which is also highlighted with a red box), 'Discord', 'Refer a Friend', and 'Log out'. At the bottom, there are links for 'Discord', 'F.A.Q.', 'About', 'Terms', 'Privacy', 'Cookies', and a light mode switch.

featherless.ai

Models Resources Pricing Chat Status

Search models...

Darin Verheijke

Account & Billing

API Keys

FEATHERLESS

API Keys

Private Models

Private Cloud

ACCOUNT

Profile

Change Password

Subscription

Payment Methods

Referral Program

API Keys

aifu

rc\_2e717\*\*\*\*\*428a

Created: Apr 3, 2025 Last Used: Oct 29, 2025

litellm

rc\_14a6a\*\*\*\*\*f182

Created: Jun 2, 2025 Last Used: Oct 16, 2025

© 2025 Featherless.ai

Discord F.A.Q. About Terms Privacy Cookies

# Get Started in 3 Steps

## Choose a model

- Go to our [model catalog](#)
- Pick a model for your use case and copy the model ID
- Test with a simple chat completion call!

The screenshot shows the featherless.ai website's model catalog page. The top navigation bar includes links for Models (which is highlighted with a red box and a red arrow pointing to it), Resources, Pricing, Chat, and Status. A search bar and a user profile icon are also present. The main content area is titled "Models" and displays a grid of 12 model cards. Each card contains the model name, a "Warm" status indicator, a brief description, a thumbnail image, a like count, and a cost. The models shown include various families like Deepseek 3, Gemma, GLM, GPT OSS, GPT-SW3, GPT2-SW3, Ling 2, Llama 3.3, Llama 3.2, Llama 3.1, Llama 3, Llama 2, Mellum, Mistral 3.1, and Mistral 3. The left sidebar lists "MODEL FAMILY" categories with their counts.

MODEL FAMILY	COUNT
Deepseek 3	(4)
Gemma 3	(179)
Gemma 2	(840)
Gemma	(13)
GLM 4.6	(1)
GLM 4	(20)
GPT OSS	(1)
GPT-SW3	(5)
GPT2-SW3	(5)
Ling 2	(1)
Llama 3.3	(147)
Llama 3.2	(2505)
Llama 3.1	(1192)
Llama 3	(1188)
Llama 2	(1054)
Mellum	(6)
Mistral 3.1	(18)
Mistral 3	(120)

1 - 20 of 12167

Model	Description	Thumbnail	Warm	Like Count	Cost
qwen25-7b-lc	Qwen/Qwen2.5-7B-Instruct		Warm	852	8.137.502
qwen3-0b6	Qwen/Qwen3-0.6B		Warm	755	7.199.888
qwen25-0b5	Gensyn/Qwen2.5-0.5B-Instruct		Warm	26	6.456.472
llama31-8b	meta-llama/Llama-3.1-8B-Instruct		Warm	4.863	5.172.472
llama31-8b	meta-llama/Meta-Llama-3.1-8B-Instruct		Warm	4.863	5.172.472
qwen3-4b	Qwen/Qwen3-4B-Instruct-2507		Warm	440	4.008.995
tinyllama-1b1	TinyLlama/TinyLlama-1.1B-Chat-v1.0		Warm	1.439	3.955.967
llama32-1b	meta-llama/Llama-3.2-1B-Instruct		Warm	1.143	3.814.050
qwen25-3b	Qwen/Qwen2.5-3B-Instruct		Warm	325	3.651.458

# Popular Model Options



DeepSeek-V3-0324

Advanced reasoning and  
coding capabilities



Llama-3.3-70B-Instruct

Powerful general-purpose  
model



Kimi-K2-Instruct

Long context capabilities



Mistral-Nemo-Instruct

Fast and efficient processing



GLM-4.6

Bilingual capabilities

# Your First API Call

## Using Python Requests

```
import requests
response = requests.post(
    url="https://api.featherless.ai/v1/chat/completions",
    headers={
        "Authorization": "Bearer YOUR_API_KEY"
    },
    json={
        "model": "deepseek-ai/DeepSeek-V3-0324",
        "messages": [
            {"role": "user", "content": "Hello!"}
        ]
    }
)
print(response.json())
```

## Using OpenAI SDK

```
from openai import OpenAI
client = OpenAI(
    base_url="https://api.featherless.ai/v1",
    api_key="YOUR_API_KEY"
)
response = client.chat.completions.create(
    model="deepseek-ai/DeepSeek-V3-0324",
    messages=[
        {"role": "user", "content": "Hello!"}
    ]
)
print(response.choices[0].message.content)
```



# Core API Endpoints

## /v1/chat/completions

**Best for:** Chatbots, virtual assistants, conversational applications

- Structured messages with roles (system, user, assistant)
- Maintains conversation context automatically
- Ideal for interactive user-assistant interactions

## /v1/completions

**Best for:** Content generation, text transformation, data extraction

- Takes a single prompt string
- Direct control over prompt format
- Maximum flexibility for custom use cases

# Troubleshooting Common Errors



## 401 - Unauthenticated

API key not recognized. Verify you've copied it correctly or generate a new one from account settings.



## 403 - Unauthorized

Model is gated. Visit the model's page, click "Unlock Model," and agree to license terms.



## 500 - Internal Server Error

Request could not be processed. Check for unsupported parameters in API documentation.



## 503 - Service Unavailable

Insufficient capacity or cold model. Retry the request. If it persists after three attempts, report on Discord.

# Concurrency Limits

Featherless operates on a unique model of **capacity reservation** rather than token-based billing. Your subscription tier determines the maximum size and number of concurrent AI model inference calls you can make, ensuring consistent performance without unpredictable token costs.

## Model Concurrency Costs

7B to 15B	1	Qwen 2.5 7B, Llama2 13B
24B to 34B	2	Qwen 32B Coder, Mistral 3 24B
70B and 72B	4	Llama 3.3 70B, Qwen 2.5 72B
Deepseek v3, R1 & Kimi-K2	4	(Feather Premium only)

For a more advanced explanation and detailed examples, visit our [Concurrency Limits documentation](#).

# Application Guides

Explore our detailed application guides designed to help you seamlessly integrate Featherless.ai into your existing projects and workflows.

Find these guides and more on our documentation page:

[featherless.ai/docs/application-guides](https://featherless.ai/docs/application-guides)

## Coding

### Aider

Powering your in-terminal coding assistance with inference from featherless

### Cursor

Setting Up Cursor with Featherless.ai

### Roo Code

Roo Code (prev. Roo Cline) gives you a whole dev team of AI agents in your code editor.

### Cline

A fully collaborative AI partner that's open source, fully extensible, and designed to amplify developer impact.

## Workflow Automation

### n8n

Supercharge your AI workflows by connecting Featherless with n8n's powerful automation platform. Build sophisticated AI-powered automations that connect with 1000+ apps and services.

### Dify

Unleash 10,000+ open-source models in Dify with our dedicated plugin. Build powerful AI applications using serverless inference and low-code orchestration without the infrastructure headaches.

# Ready to Build!

You now have everything you need to start building amazing AI-powered applications. Here's your action plan:

## 1 Test Your Setup

Get your API key and test with a simple example

## 2 Explore Models

Find the right model for your use case in the catalog

## 3 Review Examples

Check the GitHub cookbook for code samples

## 4 Start Building

Begin your hackathon project with confidence

---

Quick Links: [Sign Up](#) · [Model Catalog](#) · [Documentation](#) · [GitHub](#)

Good luck with your hackathon project! We can't wait to see what you build!